

1: R Programming Matrices and Data Frames
2: No. Let me start something new.

Selection: 1

| Let's create a vector containing the numbers 1 through 20 using the `:` operator. Store the result in a variable called
| my_vector.

```
> 1:20  
[1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20
```

| Not quite, but you're learning! Try again. Or, type info() for more options
| .

| You learned about the `:` operator in the lesson on sequences. If you wanted to create a vector containing the numbers 1,
| 2, and 3 (in that order), you could use either c(1, 2, 3) or 1:3. In this case, we want the numbers 1 through 20 stored in
| a variable called my_vector. Also, remember that you don't need the c() function when using `:`.

```
> my_vector <- 1:20
```

| Nice work!

```
| =====  
| 9%
```

| View the contents of the vector you just created.

```
>  
> my_vector  
[1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20
```

| That's a job well done!

```
| =====  
| 11%
```

| The dim() function tells us the 'dimensions' of an object. What happens if we do dim(my_vector)? Give it try.

```
> dim(my_vector)  
NULL
```

| All that hard work is paying off!

```
| =====  
| 14%
```

| Clearly, that's not very helpful! Since my_vector is a vector, it doesn't have a `dim` attribute (so it's just NULL), but
| we can find its length using the length() function. Try that now.

```
> length(my_vector)  
[1] 20
```

| That's the answer I was looking for.

```
| =====  
| 17%
```

| Ah! That's what we wanted. But, what happens if we give my_vector a `dim` attribute? Let's give it a try. Type
| `dim(my_vector) <- c(4, 5)`.

```
> dim(my_vector) <- c(4,5)
```

| Excellent job!

|=====

| It's okay if that last command seemed a little strange to you. It should! The `dim()` function allows you to get OR set the
| `dim` attribute for an R object. In this case, we assigned the value `c(4, 5)` to the `dim` attribute of `my_vector`.

...

|=====

| Use `dim(my_vector)` to confirm that we've set the `dim` attribute correctly.

```
> dim(my_vector)
[1] 4 5
```

| Great job!

|=====

| Another way to see this is by calling the `attributes()` function on `my_vector`. Try it now.

```
> attributes(my_vector)
$dim
[1] 4 5
```

| Nice work!

|=====

| Just like in math class, when dealing with a 2-dimensional object (think rectangular table), the first number is the
| number of rows and the second is the number of columns. Therefore, we just gave `my_vector` 4 rows and 5 columns.

...

|=====

| But, wait! That doesn't sound like a vector any more. Well, it's not. Now it's a matrix. View the contents of `my_vector`
| now to see what it looks like.

```
>
> my_vector
  [,1] [,2] [,3] [,4] [,5]
[1,]   1   5   9  13  17
[2,]   2   6  10  14  18
```

```
[3,] 3 7 11 15 19
[4,] 4 8 12 16 20
```

| Excellent job!

```
|=====
| 34%
```

| Now, let's confirm it's actually a matrix by using the class() function. Type class(my_vector) to see what I mean.

```
> class(my_vector)
[1] "matrix"
```

| That's the answer I was looking for.

```
|=====
| 37%
```

| Sure enough, my_vector is now a matrix. We should store it in a new variable that helps us remember what it is. Store the value of my_vector in a new variable called my_matrix.

```
> my_matrix <- my
Error: object 'my' not found
> my_matrix <- my_vector
```

| You are amazing!

```
|=====
| 40%
```

| The example that we've used so far was meant to illustrate the point that a matrix is simply an atomic vector with a dimension attribute. A more direct method of creating the same matrix uses the matrix() function.

...

```
|=====
| 43%
```

| Bring up the help file for the matrix() function now using the `?` function.

```
> ?matrix()
```

| You're close...I can feel it! Try it again. Or, type info() for more options.

| The command ?matrix will do the trick.

```
> ?matrix
```

| You're the best!

```
|=====
| 46%
```

| Now, look at the documentation for the matrix function and see if you can figure out how to create a matrix containing the same numbers (1-20) and dimensions (4 rows, 5 columns) by calling the matrix() function. Store the result in a variable called my_matrix2.

```
> my_matrix2 <- matrix(1:20, nrow = 4, ncol = 5)
```

| You are doing so well!

```
|=====
| 49%
```

| Finally, let's confirm that my_matrix and my_matrix2 are actually identical. The identical() function will tell us if its first two arguments are the same. Try it out.

```
> identical(my_matrix, my_matrix2)
[1] TRUE
```

| Keep up the great work!

```
|=====
| 51%
```

| Now, imagine that the numbers in our table represent some measurements from a clinical experiment, where each row represents one patient and each column represents one variable for which measurements were taken.

...

```
|=====
| 54%
```

| We may want to label the rows, so that we know which numbers belong to each patient in the experiment. One way to do this is to add a column to the matrix, which contains the names of all four people.

...

```
|=====
| 57%
```

| Let's start by creating a character vector containing the names of our patients -- Bill, Gina, Kelly, and Sean. Remember that double quotes tell R that something is a character string. Store the result in a variable called patients.

```
>
> c("Bill","Gina","kelly","Sean")
[1] "Bill" "Gina" "kelly" "Sean"
```

| That's not exactly what I'm looking for. Try again. Or, type info() for more options.

| Make sure to capitalize the first letter of each name and to store the result in a variable called patients. Also, don't get the order of the patients mixed up! That would be a disaster!

```
> patients <- c("Bill","Gina","kelly","Sean")
```

| You got it!

```
|=====
| 60%
```

| Now we'll use the `cbind()` function to 'combine columns'. Don't worry about storing the result in a new variable. Just call
| `cbind()` with two arguments -- the `patients` vector and `my_matrix`.

```
> cbind(my_matrix, patients)
      patients
[1,] "1" "5" "9" "13" "17" "Bill"
[2,] "2" "6" "10" "14" "18" "Gina"
[3,] "3" "7" "11" "15" "19" "Kelly"
[4,] "4" "8" "12" "16" "20" "Sean"
```

| Give it another try. Or, type `info()` for more options.

| Type `cbind(patients, my_matrix)` to add the names of our patients to the matrix of numbers.

```
> cbind(patients, my_matrix)
      patients
[1,] "Bill" "1" "5" "9" "13" "17"
[2,] "Gina" "2" "6" "10" "14" "18"
[3,] "Kelly" "3" "7" "11" "15" "19"
[4,] "Sean" "4" "8" "12" "16" "20"
```

| Perseverance, that's the answer.

```
| =====
| 63%
```

| Something is fishy about our result! It appears that combining the character vector with our matrix of numbers caused
| everything to be enclosed in double quotes. This means we're left with a matrix of character strings, which is no good.

...

```
| =====
== | 66%
```

| If you remember back to the beginning of this lesson, I told you that matrices can only contain ONE class of data.
| Therefore, when we tried to combine a character vector with a numeric matrix, R was forced to 'coerce' the numbers to
| characters, hence the double quotes.

...

```
| =====
==== | 69%
```

| This is called 'implicit coercion', because we didn't ask for it. It just happened. But why didn't R just convert the
| names of our patients to numbers? I'll let you ponder that question on your own.

...

```
| =====
===== | 71%
```

| So, we're still left with the question of how to include the names of our patients in the table without destroying the
| integrity of our numeric data. Try the following -- `my_data <- data.frame(patients, my_matrix)`

```
> my_data <- data.frame(patients,my_matrix)
```

| Excellent job!

```
|=====|
=====| 74%
```

| Now view the contents of my_data to see what we've come up with.

```
> my_data
  patients x1 x2 x3 x4 x5
1    Bill  1  5  9 13 17
2     Gina  2  6 10 14 18
3    kelly  3  7 11 15 19
4     Sean  4  8 12 16 20
```

| Nice work!

```
|=====|
=====| 77%
```

| It looks like the data.frame() function allowed us to store our character vector of names right alongside our matrix of numbers. That's exactly what we were hoping for!

...

```
|=====|
=====| 80%
```

| Behind the scenes, the data.frame() function takes any number of arguments and returns a single object of class `data.frame` that is composed of the original objects.

...

```
|=====|
=====| 83%
```

| Let's confirm this by calling the class() function on our newly created data frame.

```
>
> class(my_data)
[1] "data.frame"
```

| Keep working like that and you'll get there!

```
|=====|
=====| 86%
```

| It's also possible to assign names to the individual rows and columns of a data frame, which presents another possible way of determining which row of values in our table belongs to each patient.

...

```
|=====|
=====| 89%
```

| However, since we've already solved that problem, let's solve a different problem by assigning names to the columns of our data frame so that we know what type of measurement each column represents.

...

```
|=====
=====| 91%
```

| Since we have six columns (including patient names), we'll need to first create a vector containing one element for each column. Create a character vector called `cnames` that contains the following values (in order) -- "patient", "age", "weight", "bp", "rating", "test".

```
> cnames <- c("patient","age","weight","bp","rating","test")
```

| You are really on a roll!

```
|=====
=====| 94%
```

| Now, use the `colnames()` function to set the ``colnames`` attribute for our data frame. This is similar to the way we used the `dim()` function earlier in this lesson.

```
> colnames(my_data)
[1] "patients" "x1"      "x2"      "x3"      "x4"      "x5"
```

| Not quite! Try again. Or, type `info()` for more options.

| Try `colnames(my_data) <- cnames`.

```
> colnames(my_data) <- cnames
```

| Great job!

```
|=====
=====| 97%
```

| Let's see if that got the job done. Print the contents of `my_data`.

```
> my_data
  patient age weight bp rating test
1   Bill   1     5   9     13    17
2   Gina   2     6  10     14    18
3  Kelly   3     7  11     15    19
4   Sean   4     8  12     16    20
```

| You're the best!

```
|=====
=====| 100%
```

| In this lesson, you learned the basics of working with two very important and common data structures -- matrices and data frames. There's much more to learn and we'll be covering more advanced topics, particularly with respect to data frames, in future lessons.